

Row-major order and column-major order

widely-used strategies to store multidimensional arrays in linear storage

0	1	2
3	4	5

Row-major order

0	1	2	3	4	5
---	---	---	---	---	---

Lexicographic order

A_{00}	A_{01}	A_{02}	A_{10}	A_{11}	A_{12}
----------	----------	----------	----------	----------	----------

A_{00}	A_{01}	A_{02}
A_{10}	A_{11}	A_{12}

Column-major order

0	3	1	4	2	5
---	---	---	---	---	---

Colexicographic order

A_{00}	A_{10}	A_{01}	A_{11}	A_{02}	A_{12}
----------	----------	----------	----------	----------	----------

Note: zero-indexing

E. W. Dijkstra, "Why numbering should start at zero" (1982)

Why is storage order important?

Caching

CPU

Cache

Small and fast

Main memory

Large and slow

0	1	2
3	4	5

3	4	5
---	---	---

printout of values: 0 1 2 3 4 5

(spatial locality)

cache misses: 2

0	1	2	3	4	5
---	---	---	---	---	---

Row-major order

CPU

Cache

Small and fast

Main memory

Large and slow

0	1	2
3	4	5

4	2	5
---	---	---


printout of values: 0 1 2 3 4 5

cache misses: 4

0	3	1	4	2	5
---	---	---	---	---	---

Column-major order

Memory address formulas



Row-major order:



index = 3


$$\text{index} = i \cdot \text{cols} + j$$

$$i = 1$$

$$j = 0$$

$$\text{index} = 1 \cdot 3 + 0$$

Column-major order:



index = 1

$$\text{index} = j \cdot \text{rows} + i$$

$$i = 1$$

$$j = 0$$

Higher dimensions

Suppose a D -dimensional tensor with shape $N_0 \times N_1 \times \dots \times N_{D-1}$

dim d is indexed by i_d

$i_d \in \{0, \dots, N_d - 1\}$

$d \in \{0, \dots, D - 1\}$

Row-major order:

$$\text{index} = i_{d-1} + N_{d-1} \cdot (i_{d-2} + N_{d-2} \cdot (\dots + N_1 i_0))$$

↑
Last dimension is contiguous ("moves fastest")

Column-major order:

$$\text{index} = i_0 + N_0 \cdot (i_1 + N_1 \cdot (\dots + N_{d-2} i_{d-1}))$$

↑
Zeroth dimension is contiguous

Conventions in different languages

Row-major order:

C

C++

Pascal

ML Libraries

PyTorch

TensorFlow

Column-major order:

FORTRAN

R

MATLAB

Julia